

Advanced Hbase

Student Name: Subbham Vishal

Course: Big Data Hadoop & Spark Training

Assignment –
To-do

Case Study Description

Let us take up the CUSTOMER and TRANSACTIONS table we have created in the

Let's Do Together section. Let us solve the following use cases using these tables :-

1. Find out the number of transaction done by each customer (These should be take up in module 8 itself)
2. Create a new table called TRANSACTIONS_COUNT. This table should have 3 fields - custid, fname and count. (Again to be done in module 8)
3. Now write a hive query in such a way that the query populates the data obtained in Step 1 above and populate the table in step 2 above. (This has to be done in module 9).
4. Now lets make the TRANSACTIONS_COUNT table Hbase complaint. In the sence, use Ser Des And Storate handler features of hive to change the TRANSACTIONS_COUNT table to be able to create a TRANSACTIONS table in Hbase. (This has to be done in module 10)
5. Now insert the data in TRANSACTIONS_COUNT table using the query in step 3 again, this should populate the Hbase TRANSACTIONS table automatically (This has to be done in module 10)
6. Now from the Hbase level, write the Hbase java API code to access and scan the TRANSACTIONS table data from java level.

Problem Solution

Case Study 2 – Customers and Transactions Data

The required tables CUSTOMERS and TRANSACTIONS are created using the To-do session of Hive basics session.

Creating the tables in hive.

```
hive> CREATE TABLE CUSTOMER(  
  > custid INT,  
  > fname STRING,  
  > lname STRING,  
  > age INT,  
  > profession STRING)  
  > row format delimited fields terminated by ',';  
OK  
Time taken: 20.704 seconds  
hive> LOAD DATA LOCAL INPATH "/home/acadgild/user_acadgild/custs.txt" into table CUSTOMER  
  > ;  
Loading data to table default.customer  
OK  
Time taken: 5.528 seconds  
  
hive> CREATE TABLE TRANSACTIONS(  
  > txnno INT,  
  > txndate STRING,  
  > custno INT,  
  > amount DOUBLE,  
  > category STRING,  
  > product STRING,  
  > city STRING,  
  > state STRING,  
  > spendby STRING)  
  > row format delimited fields terminated by ',';  
OK  
Time taken: 0.426 seconds  
hive> LOAD DATA LOCAL INPATH "/home/acadgild/user_acadgild/txn.txt" into table TRANSACTIONS;  
Loading data to table default.transactions  
OK  
Time taken: 1.995 seconds  
hive> █
```

The data from both these tables can be seen below.

The data from both these tables can be seen below.

```
hive> select * from customer;
OK
101      Amitabh Bacchan 65      Actor
102      Sharukh Khan   45      Doctor
103      Akshay  Kumar   38      Dentist
104      Anubahv kumar  58      Business
105      Pawan  Trivedi  34      service
106      Aamir   Null    42      scientest
107      Salman  Khan   43      Surgen
108      Ranbir  Kapoor  26      Industrialist
Time taken: 4.404 seconds, Fetched: 8 row(s)
```

```
hive> select * from transactions;
OK
97834  05/02/2018    101    965.0  Entertainment  Movie  Pune  Maharashtra  Daughter
98396  12/01/2018    102    239.0  Food  Grocery  Patna  Bihar  Self
34908  06/01/2018    101    875.0  Travel  Air  Bangalore  Karnataka  Spouse
70958  17/02/2018    104    439.0  Food  Restaurant  Delhi  Delhi  Wife
9874   21/01/2018    105    509.0  Entertainment  Park  Kolkata  West Bengal  NULL
94585  19/01/2018    106    629.0  Rent  House  Hyderabad  Telangana  Self
45509  20/01/2018    107    953.0  Travel  Rail  Chennai  Tamil Nadu  Brother
7864   01/02/2018    108    569.0  Rent  Parking  Goa  Goa  Wife
Time taken: 0.498 seconds, Fetched: 8 row(s)
```

Objective 1:

Find out the number of transaction done by each customer.

Only the custno and the number of transactions can be queried only the transactions table.

```
select custno, count(*) from TRANSACTIONS group by custno;
```

```
hive> select custno, count(*) from TRANSACTIONS group by custno;  
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different  
tion engine (i.e. spark, tez) or using Hive 1.X releases.
```

```
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU:  
Total MapReduce CPU Time Spent: 4 seconds 780 msec  
OK  
101      2  
102      1  
104      1  
105      1  
106      1  
107      1  
108      1  
Time taken: 59.681 seconds, Fetched: 7 row(s)
```

The custid, along with the customer's first name can be queried as follows.

```
select t.custno,c.fname, count(*) from CUSTOMER c, TRANSACTIONS t where  
c.custid=t.custno group by c.fname,t.custno;
```

```
MapReduce Jobs Launched:  
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU:  
Total MapReduce CPU Time Spent: 7 seconds 450 msec  
OK  
106      Aamir      1  
101      Amitabh    2  
104      Anubahv    1  
105      Pawan      1  
108      Ranbir     1  
107      Salman     1  
102      Sharukh    1  
Time taken: 136.514 seconds, Fetched: 7 row(s)
```

The custid, customer name and the count of times the customer has occurred in the Transactions table is shown, thus showing the number of transactions per each customer.

Objective 2:

Create a new table called **TRANSACTIONS_COUNT**. This table should have 3 fields - **custid**, **fname** and **count**.

```
create table TRANSACTIONS_COUNT(  
    custid INT,  
    fname STRING,  
    count INT)  
row format delimited fields terminated by ',';
```

```
hive> create table TRANSACTIONS_COUNT(  
    > custid INT,  
    > fname STRING,  
    > count INT)  
    > row format delimited fields terminated by ',';  
OK  
Time taken: 1.202 seconds
```

The table with the given column names is created.

Objective 3:

Now write a hive query in such a way that the query populates the data obtained in Step 1 above and populate the table in step 2 above.

```
insert overwrite table TRANSACTIONS_COUNT select t.custno,c.fname, count(*)  
from CUSTOMER c, TRANSACTIONS t where c.custid=t.custno group by  
c.fname,t.custno;
```

```
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 7.77 sec  
Total MapReduce CPU Time Spent: 7 seconds 770 msec  
OK  
Time taken: 125.207 seconds  
hive> select * from TRANSACTIONS_COUNT;  
OK  
106      Aamir      1  
101      Amitabh   2  
104      Anubahv   1  
105      Pawan     1  
108      Ranbir    1  
107      Salman    1  
102      Sharukh   1  
Time taken: 0.54 seconds, Fetched: 7 row(s)
```

The data from step 1 is inserted into the newly created table **TRANSACTIONS_COUNT**. The first line OK in the above screenshot is from when the command is run. And when queried the new table, the data is there.

Objective 4:

Now let's make the `TRANSACTIONS_COUNT` table Hbase compliant. In the sense, use Ser Des and Storage handler features of hive to change the `TRANSACTIONS_COUNT` table to be able to create a `TRANSACTIONS` table in Hbase.

For a table to be Hbase compliant and to load data into a hbase table from hive, a table has to be created in HBASE. Then, that table name can be specified in the serde properties command in Hive.

Creation of table in Hbase.

```
create 'TRANSACTIONS','txn_details'
```

```
hbase(main):001:0> create 'TRANSACTIONS','txn_details'
0 row(s) in 11.5330 seconds
=> Hbase::Table - TRANSACTIONS
```

The table name is `TRANSACTIONS` and column family is `txn_details`. The columns from hive can be added into this column family.

Then, we have to create the Hive external table on top of HBase table that you want to populate.

```
CREATE EXTERNAL TABLE HBASE_TRANSACTIONS (custid INT, fname STRING, count INT)
  STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
  WITH SERDEPROPERTIES ("hbase.columns.mapping" = ":key,txn_details:fname,txn_details:count")
  TBLPROPERTIES("hbase.table.name"="transactions");
```

```
hive> CREATE EXTERNAL TABLE HBASE_TRANSACTIONS (custid INT, fname STRING, count INT)
> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
> WITH SERDEPROPERTIES ("hbase.columns.mapping" = ":key,txn_details:fname,txn_details:count")
> TBLPROPERTIES("hbase.table.name"="TRANSACTIONS");
OK
Time taken: 5.076 seconds
hive> █
```

An external table is created. The `HBaseStorageHandler` is used in creation of the table because the table has to be Hbase compliant.

In the serde properties, the column mappings are specified as to which column in hive table is mapped to which column in which column family in Hbase.

Then in the table properties, the name of the Hbase table is specified.

Objective 5:

Now insert the data in `TRANSACTIONS_COUNT` table using the query in step 3 again, this should populate the Hbase `TRANSACTIONS` table automatically.

Populating the newly created external table.

```
insert into HBASE_TRANSACTIONS select t.custno,c.fname, count(*) from
CUSTOMER c, TRANSACTIONS t where c.custid=t.custno group by
c.fname,t.custno;
```

```
hive> insert into HBASE_TRANSACTIONS select t.custno,c.fname, count(*) from CUSTOMER c, TRANSACTIONS t where c.custid=t.custno
group by c.fname,t.custno;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execu
tion engine (i.e. spark, tez) or using Hive 1.X releases.
```

Once the command is successfully executed, the data can be seen in Hbase table too.

```
hbase(main):003:0> scan 'TRANSACTIONS'
ROW COLUMN+CELL
101 column=txn_details:count, timestamp=1531073229277, value=2
101 column=txn_details:fname, timestamp=1531073229277, value=Amitabh
102 column=txn_details:count, timestamp=1531073229277, value=1
102 column=txn_details:fname, timestamp=1531073229277, value=Sharukh
104 column=txn_details:count, timestamp=1531073229277, value=1
104 column=txn_details:fname, timestamp=1531073229277, value=Anubahv
105 column=txn_details:count, timestamp=1531073229277, value=1
105 column=txn_details:fname, timestamp=1531073229277, value=Pawan
106 column=txn_details:count, timestamp=1531073229277, value=1
106 column=txn_details:fname, timestamp=1531073229277, value=Aamir
107 column=txn_details:count, timestamp=1531073229277, value=1
107 column=txn_details:fname, timestamp=1531073229277, value=Salman
108 column=txn_details:count, timestamp=1531073229277, value=1
108 column=txn_details:fname, timestamp=1531073229277, value=Ranbir
7 row(s) in 0.6750 seconds
```

The table contents in hbase are displayed and the 7 rows can be seen in hbase too.

So, the table in Hbase is automatically populated when the hive table is populated.

Objective 6:

Now from the Hbase level, write the Hbase java API code to access and scan the TRANSACTIONS table data from java level.

```
package com.acadgild.cs2;

import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.hbase.HBaseConfiguration;
import org.apache.hadoop.hbase.util.Bytes;
import org.apache.hadoop.hbase.client.HTable;
import org.apache.hadoop.hbase.client.Result;
import org.apache.hadoop.hbase.client.ResultScanner;
import org.apache.hadoop.hbase.client.Scan;

public class ScanTable {
    public static void main(String args[]) throws IOException {
        Configuration config = HBaseConfiguration.create();
        @SuppressWarnings({ "deprecation", "resource" })
        HTable table = new HTable(config, "TRANSACTIONS");
        Scan scan = new Scan();
        scan.addColumn(Bytes.toBytes("txn_details"), Bytes.toBytes("count"));
        scan.addColumn(Bytes.toBytes("txn_details"), Bytes.toBytes("fname"));
        ResultScanner scanner = table.getScanner(scan);
        for (Result result = scanner.next(); result != null; result = scanner.next()) {
            String Row = Bytes.toString(result.getRow());
            String name = Bytes.toString(result.getValue("txn_details".getBytes(),
                "fname".getBytes()));
            String count = Bytes.toString(result.getValue("txn_details".getBytes(),
                "count".getBytes()));
            System.out.println(Row + "," + name + "," + count);
            scanner.close();
        }
    }
}
```

```
101,Amitabh,2
102,Sharukh,1
104,Anubahv,1
105,Pawan,1
106,Aamir,1
107,Salman,1
108,Ranbir,1
```




BIG DATA DEVELOPER

ACADGILD